

CS 97SI: INTRODUCTION TO PROGRAMMING CONTESTS

Jaehyun Park

Today's Lecture

- Shortest Path Problem
- Floyd-Warshall Algorithm
- Dijkstra's Algorithm
- Bellman-Ford Algorithm
 - ▣ System of difference constraints
- Maybe: Problem Discussion

Shortest Path Problem

- Input: a weighted graph $G = (V, E)$
 - ▣ The edges can be directed or not
 - ▣ Sometimes, we allow negative edge weights
- Output: the path between two given nodes u and v that minimizes the total weight (or cost, length)
 - ▣ Sometimes, we want to compute all-pair shortest paths
 - ▣ Sometimes, we want to compute shortest paths from u to all other nodes

Floyd-Warshall Algorithm

- Given a directed weighted graph G
- Outputs a matrix D where d_{ij} is the shortest distance from node i to j
- Can detect a negative-weight cycle
- Runs in $\Theta(n^3)$ time
- Extremely easy to code
 - ▣ Coding time less than a few minutes

Floyd-Warshall Pseudocode

- Initialize D to the given cost matrix
- For $k = 1 \dots n$:
 - ▣ For all i and j :
 - $d_{ij} = \min(d_{ij}, d_{ik} + d_{kj})$
- If $d_{ij} + d_{ji} < 0$ for some i and j , then the graph has a negative weight cycle
- Done!
 - ▣ But how does this work?

How does Floyd-Warshall work?

- Define $f(i, j, k)$ as the shortest distance from i to j , using $1 \dots k$ as intermediate nodes
 - $f(i, j, n)$ is the shortest distance from i to j
 - $f(i, j, 0) = \text{cost}(i, j)$
- The optimal path for $f(i, j, k)$ may or may not have k as an intermediate node
 - If it does, $f(i, j, k) = f(i, k, k - 1) + f(k, j, k - 1)$
 - Otherwise, $f(i, j, k) = f(i, j, k - 1)$
- Therefore, $f(i, j, k)$ is the minimum of the two quantities above

How does Floyd-Warshall work?

- We have the following recurrences and base cases
 - ▣ $f(i, j, 0) = \text{cost}(i, j)$
 - ▣ $f(i, j, k) = \min\{f(i, k, k - 1) + f(k, j, k - 1), f(i, j, k - 1)\}$
- From the values of $f(\cdot, \cdot, k - 1)$, we can calculate $f(\cdot, \cdot, k)$
 - ▣ It turns out that we don't need a separate matrix for each k ; overwriting the existing values is fine
- That's how we get Floyd-Warshall algorithm

Dijkstra's Algorithm

- Given a directed weighted graph G and a source s
 - ▣ Important: The edge weights have to be nonnegative!
- Outputs a vector d where d_i is the shortest distance from s to node i
- Time complexity depends on the implementation:
 - ▣ Can be $O(n^2 + m)$, $O(m \log n)$, $O(n \log n)$
- Very similar to Prim's algorithm
- Intuition: Find the closest node to s , and then the second closest one, then the third, etc.

Dijkstra's Algorithm

- Maintain a set of nodes S , the shortest distances to which are decided
- Also maintain a vector d , the shortest distance estimate from s
- Initially, $S = \{s\}$, and $d_v = \text{cost}(s, v)$
- Repeat until $S = V$:
 - ▣ Find $v \notin S$ with the smallest d_v , and add it to S
 - ▣ For each edge $v \rightarrow u$ of cost c :
 - $d_u = \min(d_u, d_v + c)$

Bellman-Ford Algorithm

- Given a directed weighted graph G and a source s
- Outputs a vector d where d_i is the shortest distance from s to node i
- Can detect a negative-weight cycle
- Runs in $\Theta(nm)$ time
- Extremely easy to code
 - ▣ Coding time less than a few minutes

Bellman-Ford Pseudocode

- Initialize $d_s = 0$ and $d_v = \infty$ for all $v \neq s$
- For $k = 1 \dots n - 1$:
 - ▣ For each edge $u \rightarrow v$ of cost c :
 - $d_v = \min(d_v, d_u + c)$
- For each edge $u \rightarrow v$ of cost c :
 - ▣ If $d_v > d_u + c$:
 - Then the graph contains a negative-weight cycle

Why does Bellman-Ford work?

- A shortest path can have at most $n - 1$ edges
- At the k th iteration, all shortest paths of k or less edges are computed
- After $n - 1$ iterations, all distances are final: for every edge $u \rightarrow v$ of cost c , $d_v \leq d_u + c$ holds
 - ▣ Unless there is a negative-weight cycle
 - ▣ This is how the negative-weight cycle detection works

System of Difference Constraints

- Given m inequalities of the form $x_i - x_j \leq c$
- Want to find real numbers x_1, \dots, x_n that satisfy all the given inequalities
- Seemingly this has nothing to do with shortest paths
 - ▣ But it can be solved using Bellman-Ford

Graph Construction

- Create node i for every variable x_i
- Make an imaginary source node s
- Create zero-weight edges from s to all other nodes
- Rewrite the given inequalities as $x_i \leq x_j + c$
 - ▣ For each of these constraint, make an edge from j to i with weight c
- Now we run Bellman-Ford using s as the source

What happens?

- For every edge $j \rightarrow i$ with cost c , the shortest distance d vector will satisfy $d_i \leq d_j + c$
 - ▣ Setting $x_i = d_i$ gives a solution!
- What if there is a negative-weight cycle?
 - ▣ Assume that $1 \rightarrow 2 \rightarrow \dots k \rightarrow 1$ is a negative-weight cycle
 - ▣ From our construction, the given constraints contain $x_2 \leq x_1 + c_1, x_3 \leq x_2 + c_2$, etc.
 - ▣ Adding all of them gives $0 \leq$ (something negative)
 - ▣ i.e. the given constraints were impossible to satisfy

System of Difference Constraints

- It turns out that our solution minimizes the *span* of the variables: $\max x_i - \min x_i$
 - ▣ We won't prove it
 - ▣ This is a big hint on POJ 3169!